

Write a function to see if a binary tree ↴


A **binary tree** is a **tree** where every node has two or fewer children. The children are usually called left and right.

```
class BinaryTreeNode:
```

[▼ Python](#)

```
    def __init__(self, value):  
        self.value = value  
        self.left = None  
        self.right = None
```


This lets us build a structure like this:



That particular example is special because every level of the tree is completely full. There are no "gaps." We call this kind of tree "**perfect**."

Binary trees have a few interesting properties when they're perfect:

Property 1: the number of total nodes on each "level" doubles as we move down the tree.



Property 2: the number of nodes on the last level is equal to the sum of the number of nodes on all other levels (plus 1). In other words, about *half* of our nodes are on the last level.

Let's call the number of nodes n , and the height of the tree h . h can also be thought of as the "number of levels."

If we had h , how could we calculate n ?

Let's just add up the number of nodes on each level! How many nodes are on each level?

If we zero-index the levels, the number of nodes on the x th level is exactly 2^x !

1. Level **0**: 2^0 nodes,
2. Level **1**: 2^1 nodes,
3. Level **2**: 2^2 nodes,
4. Level **3**: 2^3 nodes,
5. *etc*

So our total number of nodes is:

$$n = 2^0 + 2^1 + 2^2 + 2^3 + \dots + 2^{h-1}$$

Why only up to 2^{h-1} ? Notice that we started counting our levels at 0. So if we have h levels in total, the last level is actually the " $h - 1$ "-th level. That means the number of nodes on the last level is 2^{h-1} .

But we can simplify. Property 2 tells us that the number of nodes on the last level is (1 more than) half of the total number of nodes, so we can just take the number of nodes on the last level, multiply it by 2, and subtract 1 to get the number of nodes overall. We know the number of nodes on the last level is 2^{h-1} , So:

$$n = 2^{h-1} * 2 - 1$$

$$n = 2^{h-1} * 2 - 1$$

$$n = 2^{h-1} * 2^1 - 1$$

$$n = 2^{h-1+1} - 1$$

$$n = 2^h - 1$$

So that's how we can go from ***h*** to ***n***. What about the other direction?

We need to bring the ***h*** down from the exponent. That's what logs are for!

First, some quick review. $\log_{10}(100)$ simply means, "**What power must you raise 10 to in order to get 100?**". Which is 2, because $10^2 = 100$.

We can use logs in algebra to bring variables down from exponents by exploiting the fact that we can simplify $\log_{10}(10^2)$. What power must we raise **10** to in order to get **10^2** ? That's easy—it's **2**.

So in this case we can take the \log_2 of both sides:

$$n = 2^h - 1$$

$$n + 1 = 2^h$$

$$\log_2((n + 1)) = \log_2(2^h)$$

$$\log_2(n + 1) = h$$

So that's the relationship between height and total nodes in a perfect binary tree.

is "***superbalanced***" (a new tree property we just made up).

A tree is "superbalanced" if the difference between the depths of any two leaf nodes

A **leaf node** is a tree node with no children.

It's the "end" of a path to the bottom, from the root.

is no greater than one.

Here's a sample binary tree node class:

```
class BinaryTreeNode:

    def __init__(self, value):
        self.value = value
        self.left = None
        self.right = None

    def insert_left(self, value):
        self.left = BinaryTreeNode(value)
        return self.left

    def insert_right(self, value):
        self.right = BinaryTreeNode(value)
        return self.right
```

Gotchas

Your first thought might be to write a recursive function, thinking, "the tree is balanced if the left subtree is balanced and the right subtree is balanced." This kind of approach works well for some other tree problems.

But this isn't quite true. Counterexample: suppose that from the root of our tree:

- The left subtree only has leaves at depths 10 and 11.
- The right subtree only has leaves at depths 11 and 12.

Both subtrees are balanced, but from the root we will have leaves at 3 different depths.

We could instead have our recursive function get the list of distinct leaf depths for each subtree. That could work fine. But let's come up with an iterative solution instead. It's usually better to use an iterative solution instead of a recursive one because it avoids stack overflow.

We can do this in $O(n)$ time and $O(n)$ space.

Breakdown

Sometimes it's good to start by rephrasing or "simplifying" the problem.

The requirement of "the difference between the depths of any two leaf nodes is no greater than 1" implies that we'll have to compare the depths of *all possible pairs* of leaves. That'd be expensive—if there are n leaves, there are n^2 possible pairs of leaves.

But we can simplify this requirement to require less work. For example, we could equivalently say:

- "The difference between the min leaf depth and the max leaf depth is 1 or less"
- "There are at most two distinct leaf depths, and they are at most 1 apart"

If you're having trouble with a recursive approach, try using an iterative one.

To get to our leaves and measure their depths, we'll have to traverse the tree somehow. **What methods do we know for traversing a tree?**

Depth-first and breadth-first are common ways to traverse a tree. Which one should we use here?

The worst-case time and space costs of both are the same—you could make a case for either.

But one characteristic of our algorithm is that it could **short-circuit** and return False as soon as it finds two leaves with depths more than 1 apart. So maybe we should **use a traversal that will hit leaves as quickly as possible...**

Depth-first traversal will generally hit leaves before breadth-first, so let's go with that. How could we write a depth-first walk that also keeps track of our depth?

Solution

We do a depth-first walk through our tree, keeping track of the depth as we go. When we find a leaf, we throw its depth into a list of depths if we haven't seen that depth already.

Each time we hit a leaf with a new depth, there are two ways that our tree might now be unbalanced:

1. There are more than 2 different leaf depths
2. There are exactly 2 leaf depths and they are more than 1 apart.

Why are we doing a depth-first walk and not a breadth-first one? You could make a case for either. We chose depth-first because it reaches leaves faster, which allows us to short-circuit earlier in some cases.

▼ Python

```
def is_balanced(tree_root):  
    depths = [] # we short-circuit as soon as we find more than 2  
  
    # we'll treat this list as a stack that will store tuples of (node, depth)  
    nodes = []  
    nodes.append((tree_root, 0))  
  
    while len(nodes):  
  
        # pop a node and its depth from the top of our stack  
        node, depth = nodes.pop()  
  
        # case: we found a leaf  
        if (not node.left) and (not node.right):  
  
            # we only care if it's a new depth  
            if depth not in depths:  
                depths.append(depth)  
  
            # two ways we might now have an unbalanced tree:  
            # 1) more than 2 different leaf depths  
            # 2) 2 leaf depths that are more than 1 apart  
            if (len(depths) > 2) or \  
                (len(depths) == 2 and abs(depths[0] - depths[1]) > 1):  
                return False  
  
        # case: this isn't a leaf - keep stepping down  
        else:  
            if node.left:  
                nodes.append((node.left, depth + 1))  
            if node.right:  
                nodes.append((node.right, depth + 1))  
  
    return True
```

Complexity

$O(n)$ time and $O(n)$ space.

For time, the worst case is the tree is balanced and we have to iterate over *all* n nodes to make sure.

For the space cost, we have two data structures to watch: depths and nodes.

depths will never hold more than three elements, so we can write that off as $O(1)$.

Because we're doing a depth first search, nodes will hold at most d nodes where d is the depth of the tree (the number of levels in the tree from the root node down to the lowest node). So we could say our space cost is $O(d)$.

But we can also relate d to n . In a balanced tree, d is $O(\log_2(n))$ (/concept/binary-tree#property2). And the *more unbalanced* the tree gets, the closer d gets to n .

In the worst case, the tree is a straight line of right children from the root where every node in that line also has a left child. The traversal will walk down the line of right children, adding a new left child to nodes at each step. When the traversal hits the rightmost node, nodes will hold *half* of the n total nodes in the tree. Half n is $O(n)$, so our worst case space cost is $O(n)$.

What We Learned

This is an intro to some tree basics. If this is new to you, don't worry—it can take a few questions for this stuff to come together. We have a few more coming up.

Particular things to note:

Focus on **depth-first** vs **breadth-first** traversal. You should be very comfortable with the differences between the two and the strengths and weaknesses of each.

You should also be very comfortable coding each of them up.

One tip: **Remember that breadth-first uses a queue and depth-first uses a stack** (could be the call stack or an actual stack object). That's not just a clue about implementation, it also helps with figuring out the differences in behavior. Those differences come from whether we visit

nodes in the order we see them (first in, first out) or we visit the last-seen node first (last in, first out).

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.